

SkillProof

AI-ASSISTED JUNIOR HIRING

SkillProof

“Too many applicants. Not enough proof.”

Verified junior tech hiring — from messy applications to defensible shortlists.

B2B SaaS · HR teams of 50–500 · frontend, WordPress & junior tech roles

HR DASHBOARD

A. Nguyen

4.6

Ready now

M. Costa

3.9

Ready now

J. Diallo

3.1

Trainable

P. Roy

2.2

At risk

THE PROBLEM

One junior role. Hundreds of look-alike CVs.



Marion

Head of Talent

~80-person scale-up. She posts one junior frontend role — and the inbox detonates.

“Which of these 300 people can actually do the job?”



Indistinguishable CVs

200–500+ applications per role. On paper they all look the same.



Weeks of screening noise

Manual filtering by gut feel — slow, inconsistent, hard to defend.



CV-AI learns past bias

Tools that score CVs replay who you hired before. No proof of skill.

28–35

hours of manual screening

per junior hire — before the first real interview.

Job quality → role-specific test → ranked evidence

SkillProof is the wedge between your ATS and the interview — not a replacement for either.



1

Upgrade the job ad


AI flags vague or unrealistic asks and builds a skill matrix for the role.



2

Candidates prove skills

A test generated from this job — coding + MCQ. Practice ≠ application set.



3

HR gets a ranked shortlist

Dimension scores 0–5 with cited evidence and a clear recommendation.

LIVE IN THE ROOM

Live demo

Watch the screen — four moves, ~3 minutes.

1 HR

Open a published junior role → run Check listing + skill matrix

2 CANDIDATE

Apply → build profile → take the practice test (different questions)

3 CANDIDATE

Take the application test → submit → see graded result

4 HR

Open ranked candidates → recommendation + rubric evidence



[SCREENSHOT: HR ranked results dashboard]

The memorable moment: ranked shortlist with rubric evidence.

Pre-loaded data · backup video ready

Impossible to do well — at scale — without AI



Parse messy ads at scale

Turn vague posts into a structured skill matrix and a calibrated test per role.



Grade open-ended code

Consistent rubric scoring across 60+ submissions — far beyond keyword matching.



LLM cost now viable

≈ \$0.22 / month AI COGS at 10 jobs × 30 candidates.*

Job ad



Gemini



Assessment



Grade



Rank



What stays human: HR approves ad changes, publishes the role, and makes the final hire. EU AI Act treats hiring as high-risk → human oversight by design.

* Illustrative AI cost of goods sold; not a price.

A focused wedge with software margins

TAM

\$5.25B

SAM

\$1.84B

SOM · 12–24mo

\$9.2M ARR

Wedge: junior-hiring evidence for firms of 50–500 employees.

Landscape

ATS + AI CV screen

Fast, but CV-centric — no proof of skill.

Coding test platforms

Annual lock-in; per-candidate fees bite at volume.

Manual process

Slow and not auditable.

Our edge: one workflow · practice ≠ application · explainable rubrics

Pricing — monthly subscription

Tier	Monthly	Jobs	Graded/mo
Starter (20–50)	€149	2	100
Growth (50–100) ★	€299	5	300
Pro path (250–500)	€899	20	2,000

Active jobs + graded candidates/mo · unlimited seats. Growth = most popular.

GTM: Paid pilots €500–1,500 → founder outbound (0–20) → case studies (20–60) → referrals (60–100).

We know where this breaks — and how we hold it

Market

No willingness to pay



Paid pilots up front; prove ROI in hours saved.

Tech

AI quality / latency



Guardrails, monitoring, human review on low confidence.

Adoption

Tool fatigue · ATS “good enough”



Wedge messaging; assessments linked to the job.

Regulatory

EU AI Act (Aug 2026)

Compliance

Disclosure, human oversight, documentation.



Trade-off we own: quality-at-speed — a fast first pass, then deeper review whenever the model is uncertain.

Help us prove it with real pilots



Our ask

10 paid pilots with HR leads at 50–100-employee firms — or warm intros to Heads of Talent in EU scale-ups.

Next 90 days

- ✓ Ship pilot metrics — hours saved & shortlist quality
- ✓ EU AI Act compliance checklist
- ✓ ATS export for shortlists

“We know the risks — we’re validating with real pilots and measurable ROI.”

Unit economics

80%+

Gross margin

Software economics — AI is a tiny share of COGS.

~\$0.22

AI COGS / mo*

At 10 jobs × 30 candidates on Gemini 3.1 Flash Lite.

€299

Growth ARPA

Most-popular tier; unlimited seats keeps adoption frictionless.

AI cost is effectively a rounding error against subscription revenue — margin scales with seats and volume.

** Illustrative; demo billing only, real payments not implemented.*


Funnel math

From noise to a defensible 5

250 applicants → ~60 graded → 12 shortlisted → ~5 interviews.

Every step is evidence-backed, so the final five are defensible — not gut-feel.

Architecture in one line



Grading is AI-only in the current MVP — no live code sandbox yet (a known, deliberate scope cut).

Red-team answers

Why now?

LLMs can finally parse messy job ads and grade open-ended code consistently — and cheaply enough to run per candidate.

Why not just use free GPT?

We're the workflow, not the model: job-linked tests, practice ≠ application sets, rubric evidence, and HR-grade audit trails.

Why would HR pay?

It replaces 28–35 hours of manual screening per hire with a defensible shortlist — pilots prove the ROI before they commit.